



# Advanced Algorithms

## Module 2

### Parallel Computing



# Why Parallel Computing?

- Save time, resources, memory, ...
- Who is using it?
  - Academia
  - Industry
  - Government
  - Individuals?
- Two practical motivations:
  - Application requirements
  - Architectural concerns.
- Why now?
  - Most computers including laptops are multi-core!
  - Need to therefore study how to use parallel computers.

# Conventional Wisdom in Computer Architecture



- Power Wall + Memory Wall + ILP Wall = Brick Wall
- Old CW: Uniprocessor performance 2X / 1.5 yrs
- New CW: Uniprocessor performance only 2X / 5 yrs?

# The Academic Interest

- Algorithmics and complexity
  - How to design parallel algorithms?
  - What are good theoretical models for parallel computing?
  - How to analyze parallel algorithms?
  - Can every sequential algorithm be parallelized?
  - What are some complexity classes wrt parallel computing?

# The Academic Interest

- Systems and Programming
  - How to write parallel programs?
  - What are some tools and environments.
  - How to convert algorithms to efficient implementations.
  - What are the differences to sequential programming?
  - What are the performance measures?
  - Can sequential programs be automatically converted to parallel programs?

# The Academic Interest


- Architectures
  - What are standard architectural designs?
  - What new issues are raised due to multiple cores?
  - Downstream concerns
    - Does a programmer have to worry about this?
    - How to support the systems software as architecture changes?

# The Course Coverage

- Focus on algorithms and complexity
- Models for parallel algorithms
- Algorithm design methodologies with application
  - Semi-numerical
  - Lists
  - Trees and graphs
- Complexity, characterization, and connection to sequential complexity classes.




# The PRAM Model





- The von Neumann model.

# The PRAM Model



- An extension of the von Neumann model.

# The PRAM Model

- A set of  $n$  identical processors 
- A common access shared memory
- Synchronous time steps 
- Access to the shared memory costs the same as a unit of computation.
- Different models to provide semantics for concurrent access to the shared memory
  - EREW, CREW, CRCW(Common, Arbitrary, Priority, ...)

# The Semantics

- In all cases, it is the programmer to ensure that his program meets the required semantics.
- EREW : Exclusive Read, Exclusive Write 🗨
  - No scope for memory contention.
  - Usually the weakest model, and hence algorithm design is tough.
- CREW : Concurrent Read, Exclusive Write
  - Allow processors to read simultaneously from the same memory location at the same instant.
  - Can be made practically feasible with additional hardware

# The Semantics

- **CRCW : Concurrent Read, Concurrent Write**
  - Allow processors to read/write simultaneously from/to the same memory location at the same instant.
  - Requires further specification of semantics for concurrent write. Popular variants include
    - **COMMON** : Concurrent write is allowed so long as the all the values being attempted are equal. Example: Consider finding the Boolean OR of n bits.
    - **ARBITRARY** : In case of a concurrent write, it is guaranteed that some processor succeeds and its write takes effect.
    - **PRIORITY** : Assumes that processors have numbers that can be used to decide which write succeeds.

# PRAM Model – Advantages and Drawbacks

## Advantages

- A simple model for algorithm design
- Hides architectural details for the designer.
- A good starting point


## Disadvantages

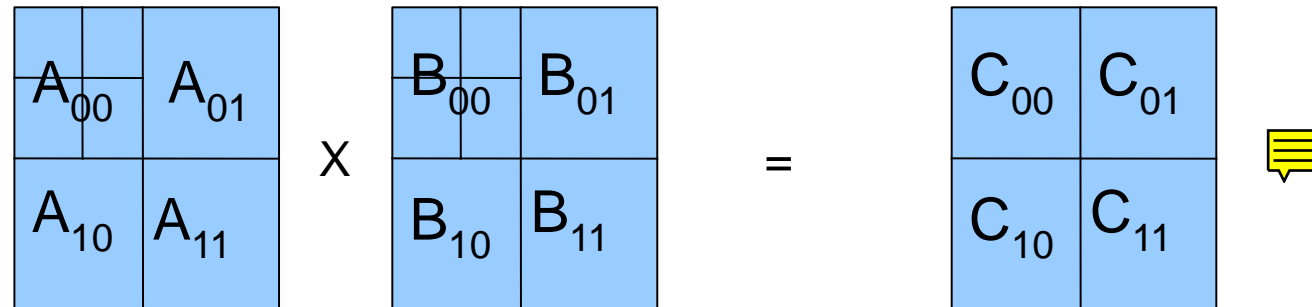
- Ignores architectural features such as:
  - memory bandwidth,
  - communication cost and latency,
  - scheduling, ...
- Hardware may be difficult to realize

# Example 1 – Matrix Multiplication

- One of the fundamental parallel processing tasks.
- Applications to several important problems in linear algebra, signal processing and optimization.
- Several techniques that work in parallel also.

# Example I – Matrix Multiplication


- Recall that in  $C = A \times B$ ,  $C[i,j] = \sum A[i,k].B[k,j]$ .
- Consider the following recursive approach: 
  - Works well in practice.


$$\begin{matrix} \begin{matrix} A_{00} & A_{01} \\ A_{10} & A_{11} \end{matrix} & \times & \begin{matrix} B_{00} & B_{01} \\ B_{10} & B_{11} \end{matrix} & = & \begin{matrix} C_{00} & C_{01} \\ C_{10} & C_{11} \end{matrix} \end{matrix}$$
$$\begin{aligned} C_{00} &= A_{00} \cdot B_{00} + A_{01} \cdot B_{10} \\ C_{01} &= A_{00} \cdot B_{01} + A_{01} \cdot B_{11} \\ C_{10} &= A_{10} \cdot B_{00} + A_{11} \cdot B_{10} \\ C_{11} &= A_{10} \cdot B_{01} + A_{11} \cdot B_{11} \end{aligned}$$



# Example I – Matrix Multiplication

- Other approaches include Cannon's algorithm



- Can overlap computation with communication.
- Works well when the number of processors is more.

# Example 2 – New Parallel Algorithm

Listing 1:

$S(1) = A(1)$

for  $i = 2$  to  $n$  do

$S(i) = S(i-1) \circ A(i)$

- **Prefix Computations:** Given an array  $A$  of  $n$  elements and an associative operation  $\circ$ , compute  $A(1) \circ A(2) \circ \dots \circ A(i)$  for each  $i$ .
- A very simple sequential algorithm exists for this problem.
- Many computations can be expressed in terms of prefix computations.


# Parallel Prefix Computation

- The sequential algorithm in Listing 1 is not efficient in parallel.
  - In particular, has to wait for the output of  $S(i)$  to compute the output  $S(i+1)$ .
- Need a new algorithm approach.
  - Balanced Binary Tree


# Balanced Binary Tree

- An **algorithm design approach** for parallel algorithms
- Many problems can be solved with this design technique.
- Easily amenable to parallelization and analysis.

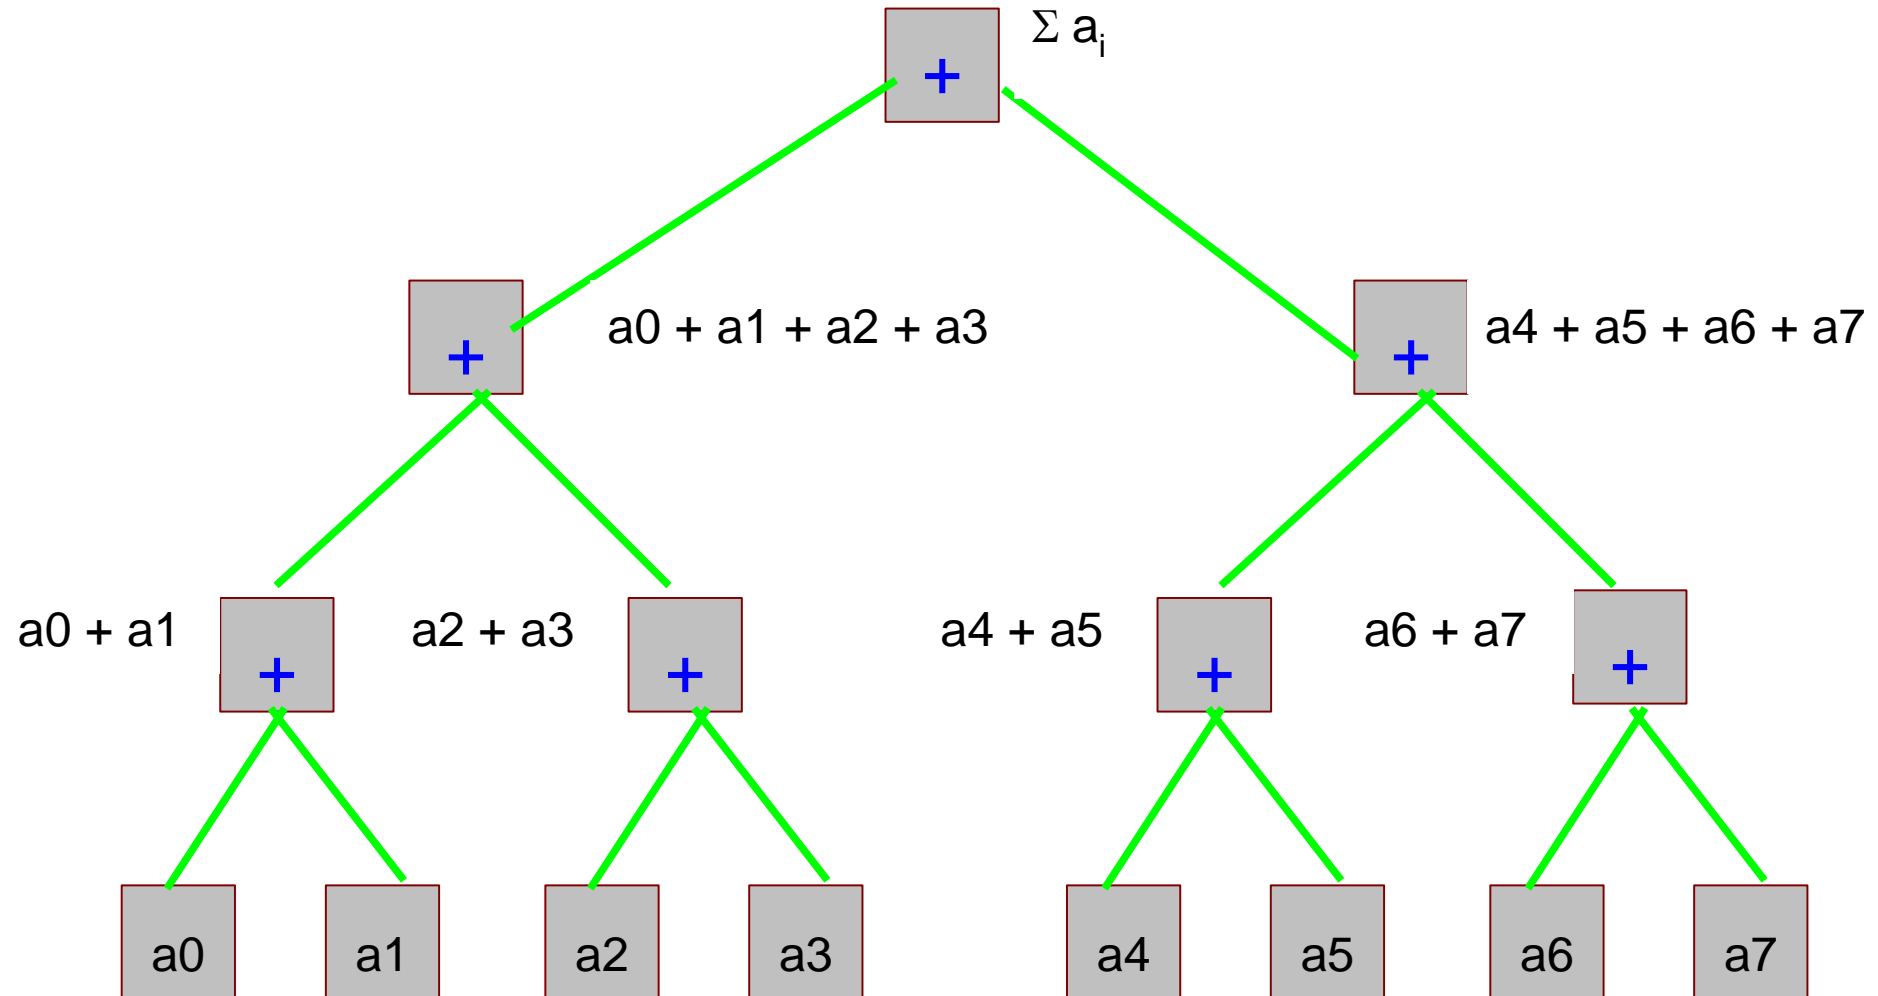
# Balanced Binary Tree

- A complete binary tree with processors at each internal node.
- Input is at the leaf nodes 
- Define operations to be executed at the internal nodes.
  - Inputs for this operation at a node are the values at the children of this node.
- Computation as a tree traversal from leaf to root.

# Balanced Binary Tree – Prefix Sums



# Balanced Binary Tree – Sum



# Balanced Binary Tree – Sum

- The above approach called as an ``upward traversal"
  - Data flow from the children to the root.
  - Helpful in other situations also such as computing the max, expression evaluation.
- Analogously, can define a downward traversal
  - Data flows from root to leaf
  - Helps in settings such as element broadcast



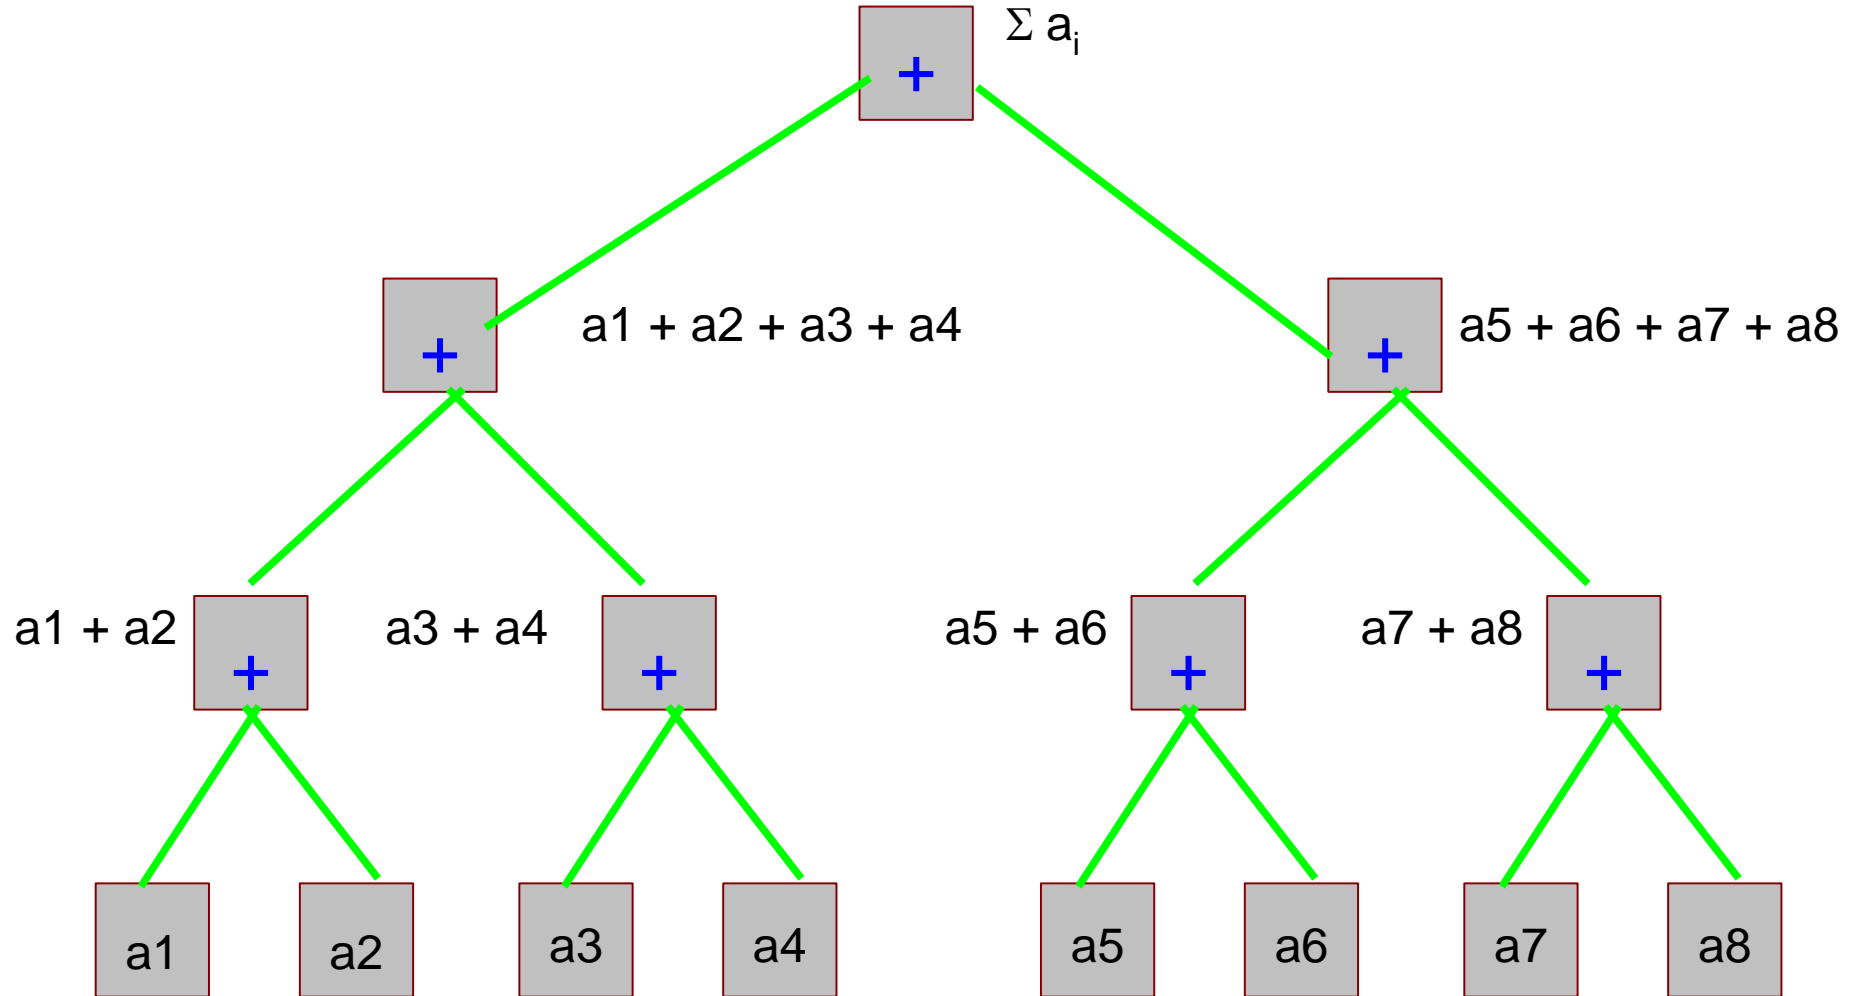
# Balanced Binary Tree

- Can use a combination of both upward and downward traversal.
- Prefix computation requires that.
- Illustration in the next slide.

A : 4 2 -1 3 0 5

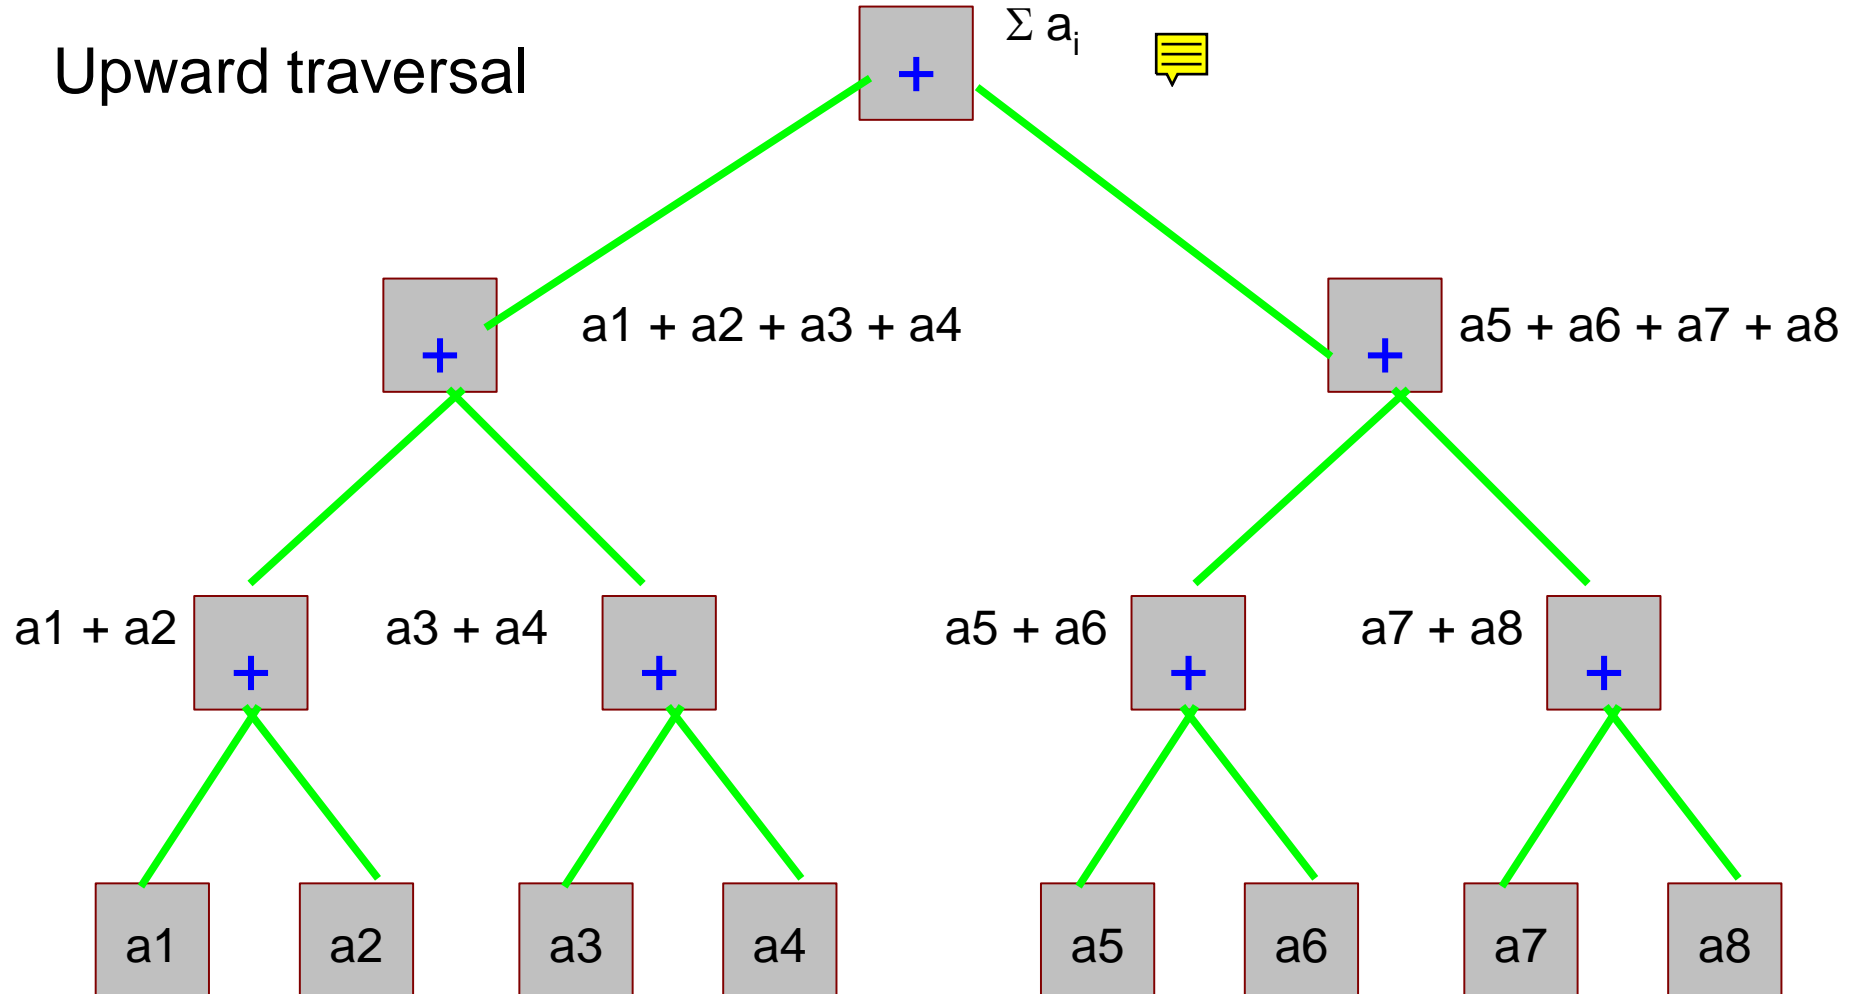
PrefixSum(A) : 4 6 5 8 8 13

# Balanced Binary Tree – Sum



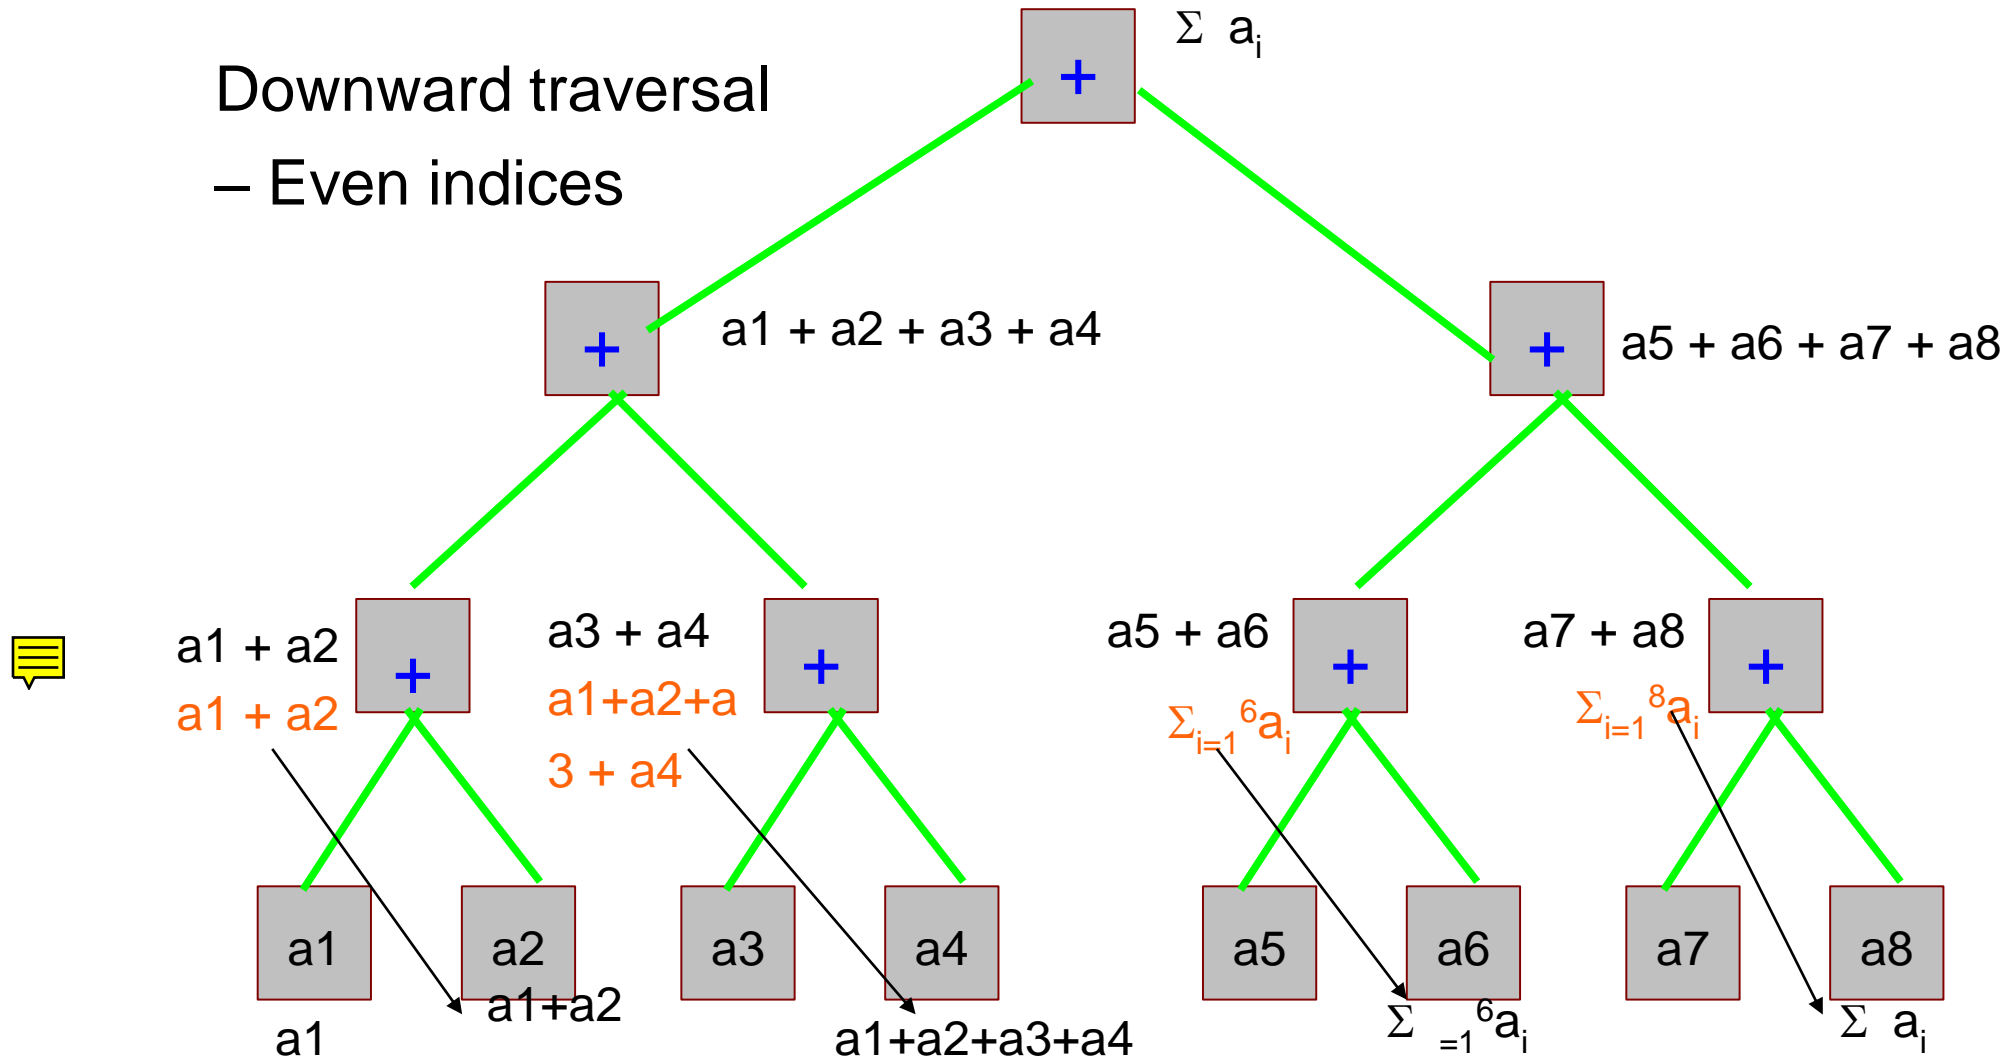
# Balanced Binary Tree – Prefix Sum

Upward traversal



# Balanced Binary Tree – Prefix Sum

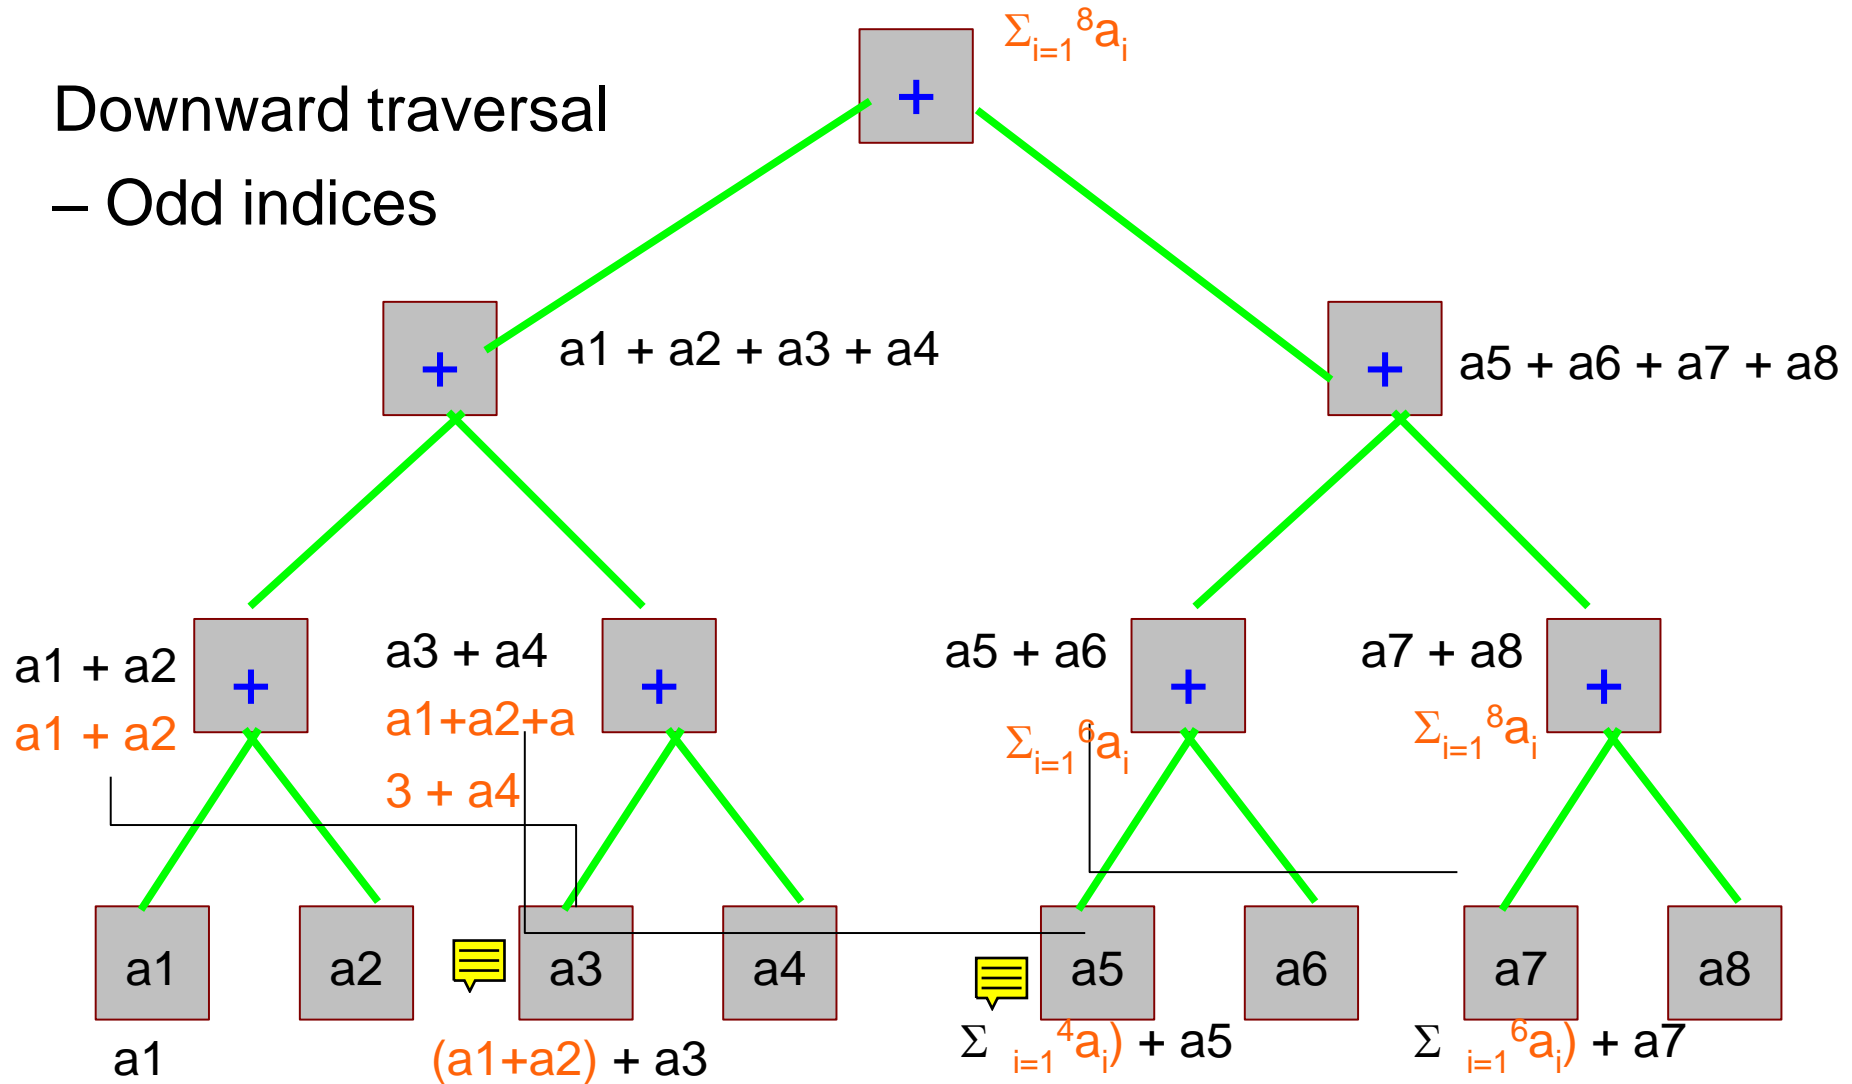
Downward traversal  
– Even indices



# Balanced Binary Tree – Prefix Sum

## Downward traversal

- Odd indices



# Balanced Binary Tree – Prefix Sums

- Two traversals of a complete binary tree.
- The tree is only a visual aid.
  - Map processors to locations in the tree
  - Perform equivalent computations.
  - Algorithm designed in the PRAM model.
  - Works in logarithmic time, and optimal number of operations.

//upward traversal

1. for  $i = 1$  to  $n/2$  do in parallel

$$b_i = a_{2i-2} \circ a_{2i}$$

2. Recursively compute the prefix sums of  $B = (b_1, b_2, \dots, b_{n/2})$  and store them in  $C = (c_1, c_2, \dots, c_{n/2})$

//downward traversal


3. for  $i = 1$  to  $n$  do in parallel

$$i \text{ is even} : s_i = c_{i/2}$$



$$i = 1 : s_1 = c_1$$

$$i \text{ is odd} : s_i = c_{(i-1)/2} \circ a_i$$

# Analysis of Parallel Algorithms



- To analyze parallel algorithms, we rely on asymptotics and recurrences.
- Each operation costs 1 unit, only sequential time needs to be counted. We assume **as many processors as can be used** are available.
- In the prefix sum example, let  $T(n)$  be the time in parallel for an input of size  $n$ .
  - Step 1 can use  $n/2$  processors in parallel each taking 1 unit of time. 
  - Step 2 is a recursive call and takes  $T(n/2)$  time.
  - Step 3 uses  $n$  processors each taking 1 unit of time.

# Analysis of Parallel Algorithms

- The recurrence relation is:
  - $T(n) = T(n/2) + O(1)$  
  - Can ignore effects due to constant factors, such as the difference in the number of processors between steps 1 and 3.
- The solution to the above recurrence is  $T(n) = O(\log n)$ .
- Another parameter of interest in parallel algorithms is the work done. 
- Can be stated as the sum of the works done by each of the processors.



# Analysis of Parallel Algorithms


- The work done by the prefix algorithm can be expressed by the recurrence
  - $W(n) = W(n/2) + O(n)$ .
  - The  $O(n)$  accounts for the work in the first and the third steps. 
  - Solution:  $W(n) = O(n)$ . 
- Work done can indicate if the algorithm is doing about the same amount of operations as the best known sequential algorithm.
- Such a parallel algorithm is called an **optimal algorithm**.




# Other Design Paradigms

- Partitioning
  - Similar to divide and conquer
  - But **no need** to combine solutions
  - Can treat problems independently and solve in parallel.
  - Example: Parallel merging, searching.

# Merging in Parallel by Partitioning



- Two sorted arrays A and B to be merged into C.
- Let A be a sorted array. Let  $\text{Rank}(x, A)$  be the number of elements smaller than x in A. 
- Claim:  $\text{Rank}(x, C) = \text{Rank}(x, A) + \text{Rank}(x, B)$
- For x in A,  $\text{Rank}(x, A)$  is immediately available. To find  $\text{Rank}(x, B)$  can use binary search in parallel.

# Quick Example


A = [8 10 12 24 ]

B = [15 17 27 32]

Element	8	10	12	24	15	17	27	32
Rank in A	0	1	2	3	3	3	4	4
Rank in B	0	0	0	2	0	1	2	3
Rank in C	0	1	2	5	3	4	6	7

C = [ 8 10 12 15 17 24 27 32 ]

# Merging in Parallel by Partitioning

- Time for each binary search is  $O(\log n)$
- Total time for merging =  $O(\log n)$ , the total work is  $O(n \log n)$ . 
  - Not work optimal as compared to the best possible sequential time complexity of  $O(n)$ .
- Can **reduce** the total work to  $O(n)$ .
  - Induce equal-sized partitions in the arrays
  - Rank one element, say the first element, from each partition
  - Use these ranks to find the ranks of the other elements, sequentially.


# An Improved Optimal Algorithm

- **General technique**
  - Solve a smaller problem in parallel
  - Extend the solution to the entire problem.
- For the first step, the problem size to be solved is guided by the factor of non-optimality of an existing parallel algorithm.

# An Improved Parallel Algorithm


- Our simple parallel algorithm is away from work optimality by a factor of  $O(\log n)$ .
- So, we should solve a problem of size  $O(n/\log n)$ .
- For this purpose, we pick every  $\log n^{\text{th}}$  element of  $A$ , and similarly in  $B$ .
- Use the simple parallel algorithm on these elements of  $A$  and  $B$ .
  - Binary search however in the entire  $A$  and  $B$ .

# An Improved Parallel Algorithm






- Let  $A_1, A_2, \dots, A_{n/\log n}$  be the elements of  $A$  ranked in  $B$ .
- These ranks induce partitions in  $B$ .
  - Define  $[B_{r(i)} \dots B_{r(i+1)}]$  as the portion of  $B$  so that  $[A(i) \dots A(i+1)]$  have ranks in.
- Can therefore merge  $[A(i) \dots A(i+1)]$  with  $[B_{r(i)} \dots B_{r(i+1)}]$  sequentially.





# An Improved Parallel Algorithm

- Such sequential merges can happen in parallel, at each index of  $A[i]$ .
- Time taken for the sequential merge is  $O(\log n + B_{r(i+1)} - B_{r(i)})$ . 
- Time:
  - Binary search:  $O(\log n)$ , with  $n/\log n$  processors. 
  - Sequential merge:  $O(\log n)$ , subject to certain conditions.   
There are also  $n/\log n$  such merges in parallel.
- Work:
  - There are  $n/\log n$  binary searches in parallel. Work =  $O(n)$ .
  - For the sequential merges too, work =  $O(n)$ .

